

Certification options

Agentic AI development — worked example

Run 1 · 2026-08-08 · scored against a scorecard frozen before research

Sit Meridian. It wins on the two criteria the spec weighted heaviest, and the best-known credential here places last.

The answer

Sit the Meridian Agentic Systems Practitioner. It is the only credential surveyed that scores above 0.7 on both criteria this spec weighted heaviest — vendor neutrality and MCP coverage — while staying inside both the money and the hours.

Then consider Tessera as a warm-up, not as the answer. It is free, open and the most protocol-led of the four, which is why it ranks second. It is also a foundations credential that does not assess delegation at all, so it satisfies the spec’s topic list only in part. Sitting it costs nothing but ten hours and would close the context-management gap before the Meridian attempt.

Do not sit Northgate. It is the best-known credential here and it finishes last. That is not a mistake in the arithmetic; it is the spec working.

Why

The ranking turns almost entirely on one weight. Recognition carries 0.05, because the stated audience is existing clients who have already seen the work, not anonymous recruiters filtering a stack of CVs. Northgate scores 0.90 on recognition and it buys almost nothing. Against that, vendor neutrality carries 0.22 — the heaviest criterion on the card — and Northgate scores 0.25, because six of its nine objectives are its own console and SDK.

Reweight for a different audience and the answer changes honestly. If the purpose were passing an automated CV filter, recognition would carry 0.25 and Northgate would win. The scorecard is not claiming Northgate is a weak credential. It is claiming Northgate is a weak credential **for this person, for this reason**, and that distinction is the whole point of writing the spec first.

Helios finishes third on a thinner evidence base than the others. Its blueprint is gated behind a training purchase, so two of its scores come from prep-provider descriptions rather than the vendor, and one criterion could not be established at all. Its position is the least stable of the four.

Ranking

#	Credential	Vendor	All-in	Score
1	Meridian Agentic Systems Practitioner	Meridian Institute	USD 195	0.80
2	Tessera Agent Foundations	Tessera Labs	USD 0	0.68
3	Helios Applied LLM Developer	Helios AI	USD 250	0.59
4	Northgate Certified Agent Engineer	Northgate Cloud	USD 340	0.57

Score is the weighted mean over the criteria scored for that candidate, on a 0–1 scale. Weights are in the scorecard, frozen before research began.

How this was decided

The scorecard below was derived from the profile — subject, benchmark, learning preferences, objectives and budget — and frozen before any credential was looked up. Research scored against it without amending it.

Criterion	Kind	Weight	What a 1.0 looks like
Vendor neutrality	quality	0.22	Blueprint teaches protocols and patterns that transfer to any toolchain
MCP coverage	coverage	0.15	MCP named as an assessed objective with a stated blueprint weight
Agent architecture	coverage	0.15	Agent loops, planning and termination assessed as their own domain
Sub-agent delegation	coverage	0.10	Multi-agent coordination assessed with named task statements
Context management	coverage	0.08	Context windows, memory and state assessed explicitly
Accessibility	constraint	0.10	Open registration, no prerequisite, sittable from anywhere
Cost fit	constraint	0.08	All-in cost well under the comfortable figure in budget.md
Time fit	constraint	0.07	Honest prep estimate under half the total study budget
Recognition	quality	0.05	Named in job postings and tenders in the target market

Hard filters

Requirements a credential either meets or does not. Failing one excludes rather than scores down, because a hard requirement that can be outweighed is not a hard requirement.

Requirement	Rule applied
Sittable remotely	Excluded if the only route is an in-person test centre
No organisational gate	Excluded if eligibility requires partner-network membership or a minimum employer size
Fits the window	Excluded if the exam is not offered before the date in objectives.md

Comparison

Criterion	Wt	Meridian Agentic Systems Practitioner	Tessera Agent Foundations	Helios Applied LLM Developer	Northgate Certified Agent Engineer
Vendor neutrality	0.22	0.85 V	0.95 V	0.45 V	0.25 V
MCP coverage	0.15	0.90 V	0.65 V	0.70 V	0.55 V
Agent architecture	0.15	0.80 V	0.45 V	0.55 S	0.85 V
Sub-agent delegation	0.10	0.70 V	0.20 V	0.40 S	0.80 V
Context management	0.08	0.60 S	0.35 V	—	0.75 V
Accessibility	0.10	1.00 V	1.00 V	0.90 V	0.70 V
Cost fit	0.08	0.85 V	1.00 V	0.60 V	0.35 V
Time fit	0.07	0.70 I	0.95 V	0.60 I	0.30 S
Recognition	0.05	0.40 S	0.15 I	0.65 S	0.90 V
Weighted total		0.80	0.68	0.59	0.57

v vendor-confirmed · **s** secondary source · **i** inferred · — not established. A high score carrying **i** is a hypothesis, not a finding.

Unscored criteria. These were not established from the public record and are excluded from the totals rather than counted as zero:

- Helios Applied LLM Developer — Context management

Excluded

Surveyed and ruled out on a hard filter. Listed because “why not that one” is a question that otherwise gets asked again on every re-run.

Credential	Vendor	Failed
Cascade Agent Architect	Cascade Systems	No organisational gate — partner tier requires 10+ employees
Brightline AI Professional	Brightline	Sittable remotely — in-person test centre only

Where to invest

Against the Meridian blueprint, the benchmark is strong in three of five assessed domains and thin in two. Both gaps are narrow and closable by building rather than reading, which is the useful kind of gap.

- **Context and state management.** The weakest scored area in the benchmark and a named Meridian objective. Build one agent that persists state across sessions and handles a context window it genuinely overflows — the overflow is the part the exam tests and the part that never comes up when the window is large enough.
- **Evaluation and error analysis.** No test directories anywhere in the evidence reviewed. Meridian assesses evaluation as a named objective, and this is also the gap most likely to matter after the exam. Add a scored evaluation harness to one existing agent; the discipline transfers whether or not the credential is ever sat.
- **Leave delegation alone.** Already strong on the evidence, and Meridian’s treatment of it is lighter than the existing work. Reading up here would feel productive and buy nothing.

Limits and what would change this

Two scores are inferred rather than confirmed. Meridian’s time fit and Tessera’s recognition carry **i** tags. Neither is load-bearing for the recommendation — Meridian still wins if both are wrong in the least favourable direction — but the gap between Meridian and Tessera would narrow.

One criterion could not be established for Helios. Its context-management coverage is not in any public document; the blueprint is gated. Helios is scored over the remaining criteria rather than penalised, so its position is the least comparable of the four.

What would change the answer. A shift in audience — if the purpose became satisfying a tender that names a specific credential, recognition would dominate and Northgate would win outright. A published Meridian blueprint revision that reduced the MCP domain below about a tenth would also move it. Both are worth a re-run rather than a guess; the intake is already on file, so the second pass is the market sweep only.

Generated by [Spec-Led-Certification](#). Scores and weights read from `scorecard/criteria.csv` and `research/scores.csv` at compile time; sources logged in `research/sources.md`.